

Machine Learning

June 15, 2025

1 Gesture Recognition using Machine Learning

Below we combine all required preprocessing steps into a single cell—matching the Data Exploration notebook (window size = 250, train step = 125, val/test step = 250)—and then run the ML workflow: feature extraction, PCA, and model tuning.

```
[1]: from google.colab import drive  
  
_drive_path = '/content/drive'  
drive.mount(_drive_path, force_remount=False)
```

Mounted at /content/drive

```
[8]: import os  
import glob  
import random  
import matplotlib.pyplot as plt  
import numpy as np  
import pandas as pd  
  
from tensorflow.keras.utils import to_categorical  
  
from scipy import stats  
  
from sklearn.preprocessing import StandardScaler  
from sklearn.decomposition import PCA  
from sklearn.model_selection import GroupShuffleSplit, ParameterGrid, cross_validate  
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.svm import SVC  
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report  
  
from tqdm import tqdm  
  
np.random.seed(42)
```

1.0.1 Data Preparation

We load, filter, segment, split, window, stack, normalize, and one-hot-encode exactly as before.

```
[3]: root_dir = "/content/drive/MyDrive/IA_EBM/PROJECT/EMG_data_for_gestures-master"
all_txt = glob.glob(os.path.join(root_dir, "**", "*.txt"), recursive=True)

# Load & concatenate
dfs = [pd.read_csv(f, sep="\t") for f in all_txt]
df = pd.concat(dfs, ignore_index=True)

# Drop classes 0 & 7 (unmarked / extended palm)
df = df[~df['class'].isin([0, 7])].reset_index(drop=True)

# Create a segment ID each time the gesture label changes
df['segment'] = (df['class'] != df['class'].shift()).cumsum()

# Build per-segment dict
gesture_datasets = {
    f"gesture_{int(cls)}_{seg}": grp
    for (cls, seg), grp in df.groupby(['class', 'segment'])
}

# train/val/test splits by segment (60/20/20)
classes = sorted({int(k.split('_')[1]) for k in gesture_datasets})
train_segs = [] ; val_segs = [] ; test_segs = []
for cls in classes:
    keys = [k for k in gesture_datasets if int(k.split('_')[1]) == cls]
    random.shuffle(keys)
    n = len(keys)
    train_segs += keys[:int(0.6*n)]
    val_segs += keys[int(0.6*n):int(0.8*n)]
    test_segs += keys[int(0.8*n):]

# sliding windows: 250 samples, 125-step train, 250-step val/test
WINDOW_SIZE = 250
train_w, val_w, test_w = {}, {}, {}
for key, grp in gesture_datasets.items():
    sig = grp.filter(like="channel").to_numpy().T
    if key in train_segs: step = 125; target = train_w
    elif key in val_segs: step = 250; target = val_w
    else: step = 250; target = test_w

    for i in range(0, sig.shape[1]-WINDOW_SIZE+1, step):
        target[f"{key}_win_{i//step}"] = sig[:, i:i+WINDOW_SIZE].T

# stack and label
def prepare(wdict):
```

```

X = np.stack([w.T for w in wdict.values()], axis=0)
y = np.array([int(k.split('_')[1]) for k in wdict.keys()])
return X, y

X_train, y_train = prepare(train_w)
X_val, y_val = prepare(val_w)
X_test, y_test = prepare(test_w)

# normalize by train stats
mean = X_train.mean((0,1), keepdims=True)
std = X_train.std((0,1), keepdims=True) + 1e-8
for arr in (X_train, X_val, X_test):
    arr[:] = (arr - mean) / std

# one-hot encode
offset = y_train.min()
num_classes = y_train.max() - offset + 1
y_train_enc = to_categorical(y_train - offset, num_classes)
y_val_enc = to_categorical(y_val - offset, num_classes)
y_test_enc = to_categorical(y_test - offset, num_classes)

```

1.1 Feature Extraction from EMG Windows

We compute a suite of time-domain features for each fixed-length window:

- Mean, std, min, max
- Skewness, kurtosis
- RMS and energy
- Zero-crossings

```
[4]: def extract_features(window):
    """
    window: ndarray (time_steps, n_channels)
    returns: concatenated feature vector for all channels
    """
    diffs = np.diff(window, axis=0) # first-order differences
    feats = [
        np.mean(window, axis=0),      # mean per channel
        np.std(window, axis=0),       # standard deviation
        np.min(window, axis=0),       # minimum value
        np.max(window, axis=0),       # maximum value
        stats.skew(window, axis=0),   # skewness
        stats.kurtosis(window, axis=0), # kurtosis
        np.sqrt(np.mean(window**2, axis=0)), # RMS
        np.sum(window**2, axis=0),      # signal energy
        np.sum((window[:-1]*window[1:])<0, axis=0) # zero-crossings
    ]
    return np.concatenate(feats)
```

```

def make_feature_set(wdict, split_name):
    X_feat, y_lbl, groups = [], [], []
    for key, w in tqdm(wdict.items(), desc=f"Extracting features_{split_name}"):
        X_feat.append(extract_features(w)) # extract features
        cls = int(key.split('_')[1]) # get class label
        y_lbl.append(cls)
        groups.append(key.rsplit('_win_', 1)[0]) # group by original segment
    return np.vstack(X_feat), np.array(y_lbl), np.array(groups)

# Build feature sets for each split
Xtr_ft, ytr_ft, grp_tr = make_feature_set(train_w, "Train")
Xva_ft, yva_ft, grp_va = make_feature_set(val_w, "Val")
Xte_ft, yte_ft, grp_te = make_feature_set(test_w, "Test")

print("Feature shapes:", Xtr_ft.shape, Xva_ft.shape, Xte_ft.shape)

```

```

Extracting features (Train): 100%| 6457/6457 [00:16<00:00,
387.91it/s]
Extracting features (Val): 100%| 1103/1103 [00:02<00:00, 417.54it/s]
Extracting features (Test): 100%| 1112/1112 [00:04<00:00, 229.61it/s]

Feature shapes: (6457, 72) (1103, 72) (1112, 72)

```

1.2 Scaling and PCA

Standardize features then apply PCA to keep 95% of variance.

```

[5]: # Standardize features
scaler = StandardScaler().fit(Xtr_ft)
Xtr_s = scaler.transform(Xtr_ft)
Xva_s = scaler.transform(Xva_ft)
Xte_s = scaler.transform(Xte_ft)

# To avoid passing NaNs
Xtr_s = np.nan_to_num(Xtr_s, nan=0.0, posinf=0.0, neginf=0.0)
Xva_s = np.nan_to_num(Xva_s, nan=0.0, posinf=0.0, neginf=0.0)
Xte_s = np.nan_to_num(Xte_s, nan=0.0, posinf=0.0, neginf=0.0)

# Apply PCA
pca = PCA(0.95, random_state=42).fit(Xtr_s)
Xtr_p = pca.transform(Xtr_s)
Xva_p = pca.transform(Xva_s)
Xte_p = pca.transform(Xte_s)

```

```
print(f"PCA reduced {Xtr_ft.shape[1]} → {Xtr_p.shape[1]} features")
```

PCA reduced 72 → 35 features

1.3 k-Nearest Neighbors Tuning

We search over `n_neighbors` using $3 \times$ repeated, group-aware 5-fold cross-validation.

```
[9]: def tune_knn(
    X_train, y_train, groups,
    ks,
    X_test=None, y_test=None,
    n_splits=5, train_size=0.8, random_state=42
):
    """
    Tune KNN (n_neighbors in ks) with GroupShuffleSplit CV,
    retrain on all training data, and (optionally) evaluate on test.
    """
    best_k, best_val = None, -np.inf

    # group-aware splitter
    cv = GroupShuffleSplit(
        n_splits=n_splits,
        train_size=train_size,
        random_state=random_state
    )

    # search ks
    for k in tqdm(ks, desc="Tuning KNN"):
        clf = KNeighborsClassifier(n_neighbors=k)
        res = cross_validate(
            clf, X_train, y_train,
            groups=groups, cv=cv,
            return_train_score=True, n_jobs=-1
        )
        tr_m, tr_v = res['train_score'].mean(), res['train_score'].var()
        v_m, v_v = res['test_score'].mean(), res['test_score'].var()
        print(f"k={k}: train {tr_m:.3f}±{tr_v:.4f}, val {v_m:.3f}±{v_v:.4f}")

        if v_m > best_val:
            best_val, best_k = v_m, k

    print("\n→ Best k = {} (val acc = {:.3f})\n".format(best_k, best_val))

    # retrain on full train set
    best_model = KNeighborsClassifier(n_neighbors=best_k)
```

```

best_model.fit(X_train, y_train)

# final evaluation & single CM display
if X_test is not None and y_test is not None:
    y_pred = best_model.predict(X_test)
    labels = np.unique(y_test) # e.g. array([1,2,3,4,5,6])
    cm = confusion_matrix(y_test, y_pred, labels=labels)

    # plot heatmap with properly aligned ticks
    fig, ax = plt.subplots(figsize=(6,6))
    im = ax.imshow(cm, interpolation='nearest', cmap='Blues')

    # positions 0..len(labels)-1, but label them with actual class names
    tick_pos = np.arange(len(labels))
    ax.set_xticks(tick_pos)
    ax.set_xticklabels(labels)
    ax.set_yticks(tick_pos)
    ax.set_yticklabels(labels)

    ax.set_xlabel("Predicted")
    ax.set_ylabel("True")
    ax.set_title("Confusion Matrix")

    # annotate counts
    thresh = cm.max() / 2
    for i in range(len(labels)):
        for j in range(len(labels)):
            color = "white" if cm[i, j] > thresh else "black"
            ax.text(j, i, cm[i, j], ha="center", va="center", color=color)

    fig.colorbar(im, ax=ax)
    plt.tight_layout()
    plt.show()

    # classification report
    print("\nClassification Report:\n",
          classification_report(y_test, y_pred, digits=3))

return best_model

```

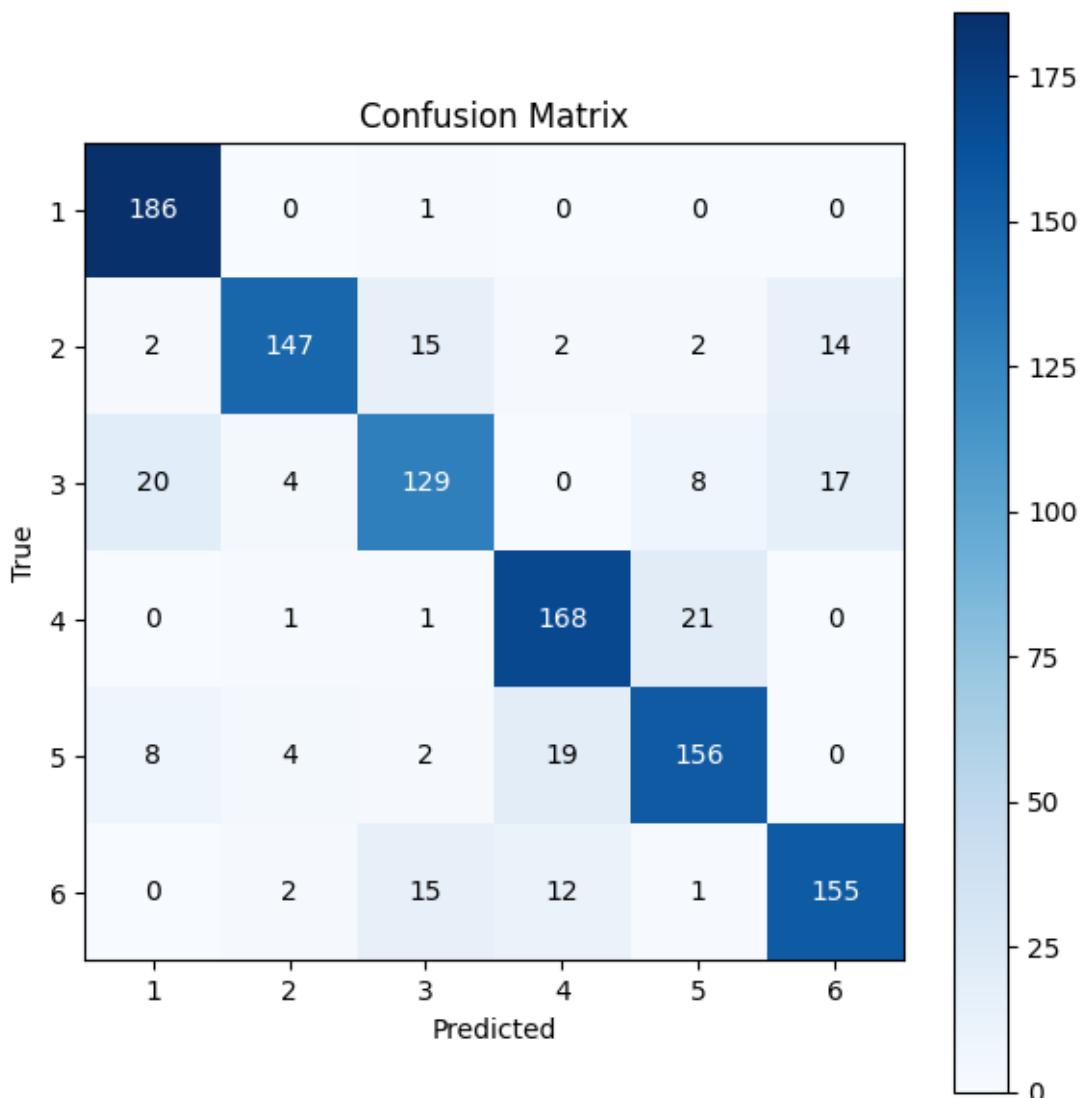
[10]: knn = tune_knn(Xtr_p, ytr_ft, grp_tr, ks=[3, 5, 7, 9, 15, 20, 50], X_test=Xte_p, ↴y_test=yte_ft)

Tuning KNN: 14% | 1/7 [00:02<00:13, 2.25s/it]

k=3: train 0.967±0.0000, val 0.790±0.0003

Tuning KNN: 29% | 2/7 [00:05<00:13, 2.63s/it]

```
k=5: train 0.948±0.0000, val 0.805±0.0003
Tuning KNN: 43%|           | 3/7 [00:09<00:14, 3.63s/it]
k=7: train 0.934±0.0000, val 0.813±0.0003
Tuning KNN: 57%|           | 4/7 [00:13<00:11, 3.70s/it]
k=9: train 0.926±0.0000, val 0.819±0.0006
Tuning KNN: 71%|           | 5/7 [00:17<00:07, 3.64s/it]
k=15: train 0.908±0.0000, val 0.820±0.0006
Tuning KNN: 86%|           | 6/7 [00:20<00:03, 3.65s/it]
k=20: train 0.897±0.0000, val 0.819±0.0007
Tuning KNN: 100%|          | 7/7 [00:23<00:00, 3.37s/it]
k=50: train 0.868±0.0000, val 0.814±0.0006
→ Best k = 15 (val acc = 0.820)
```



Classification Report:

	precision	recall	f1-score	support
1	0.861	0.995	0.923	187
2	0.930	0.808	0.865	182
3	0.791	0.725	0.757	178
4	0.836	0.880	0.857	191
5	0.830	0.825	0.828	189
6	0.833	0.838	0.836	185
accuracy			0.846	1112
macro avg	0.847	0.845	0.844	1112

weighted avg	0.847	0.846	0.845	1112
--------------	-------	-------	-------	------

1.4 Random Forest

We search over `n_estimators` and `max_depth` using repeated, group-aware CV.

```
[11]: def tune_rf(
    X_train, y_train, groups,
    param_grid,
    X_test=None, y_test=None,
    n_splits=5, train_size=0.8, random_state=42
):

    best_params, best_val = None, -np.inf

    # group-aware splitter
    cv = GroupShuffleSplit(
        n_splits=n_splits,
        train_size=train_size,
        random_state=random_state
    )

    # grid search
    for params in tqdm(ParameterGrid(param_grid), desc="Tuning RF"):
        clf = RandomForestClassifier(
            **params,
            random_state=random_state,
            n_jobs=-1
        )
        res = cross_validate(
            clf, X_train, y_train,
            groups=groups, cv=cv,
            return_train_score=True, n_jobs=-1
        )
        tr_m, tr_v = res['train_score'].mean(), res['train_score'].var()
        v_m, v_v = res['test_score'].mean(), res['test_score'].var()
        print(f"{params}: train {tr_m:.3f}±{tr_v:.4f}, val {v_m:.3f}±{v_v:.4f}")

        if v_m > best_val:
            best_val, best_params = v_m, params

    print(f"\n→ Best params = {best_params} (val acc = {best_val:.3f})\n")

    # retrain on full train set
    best_model = RandomForestClassifier(
```

```

        **best_params,
        random_state=random_state,
        n_jobs=-1
    )
best_model.fit(X_train, y_train)

# final evaluation & single CM display
if X_test is not None and y_test is not None:
    y_pred = best_model.predict(X_test)
    labels = np.unique(y_test) # e.g. array([1,2,3,4,5,6])
    cm = confusion_matrix(y_test, y_pred, labels=labels)

    # plot heatmap with properly aligned ticks
    fig, ax = plt.subplots(figsize=(6,6))
    im = ax.imshow(cm, interpolation='nearest', cmap='Blues')

    # positions 0..len(labels)-1, but label them with actual class names
    tick_pos = np.arange(len(labels))
    ax.set_xticks(tick_pos)
    ax.set_xticklabels(labels)
    ax.set_yticks(tick_pos)
    ax.set_yticklabels(labels)

    ax.set_xlabel("Predicted")
    ax.set_ylabel("True")
    ax.set_title("Confusion Matrix")

    # annotate counts
    thresh = cm.max() / 2
    for i in range(len(labels)):
        for j in range(len(labels)):
            color = "white" if cm[i, j] > thresh else "black"
            ax.text(j, i, cm[i, j], ha="center", va="center", color=color)

    fig.colorbar(im, ax=ax)
    plt.tight_layout()
    plt.show()

    # classification report
    print("\nClassification Report:\n",
          classification_report(y_test, y_pred, digits=3))

return best_model

```

[12]: rf_model = tune_rf(X_train=Xtr_p, y_train=ytr_ft, groups=grp_tr, ↴
 param_grid={'n_estimators':[50,100,200], 'max_depth':[None,5,10]}, ↴
 X_test=Xte_p, y_test=yte_ft)

```

Tuning RF: 11% | 1/9 [00:09<01:14, 9.37s/it]
{'max_depth': None, 'n_estimators': 50}: train 1.000±0.0000, val 0.829±0.0008

Tuning RF: 22% | 2/9 [00:28<01:46, 15.23s/it]
{'max_depth': None, 'n_estimators': 100}: train 1.000±0.0000, val 0.834±0.0006

Tuning RF: 33% | 3/9 [01:09<02:42, 27.07s/it]
{'max_depth': None, 'n_estimators': 200}: train 1.000±0.0000, val 0.837±0.0007

Tuning RF: 44% | 4/9 [01:14<01:30, 18.11s/it]
{'max_depth': 5, 'n_estimators': 50}: train 0.824±0.0001, val 0.788±0.0022

Tuning RF: 56% | 5/9 [01:25<01:02, 15.62s/it]
{'max_depth': 5, 'n_estimators': 100}: train 0.827±0.0000, val 0.794±0.0023

Tuning RF: 67% | 6/9 [01:56<01:02, 20.80s/it]
{'max_depth': 5, 'n_estimators': 200}: train 0.827±0.0001, val 0.795±0.0025

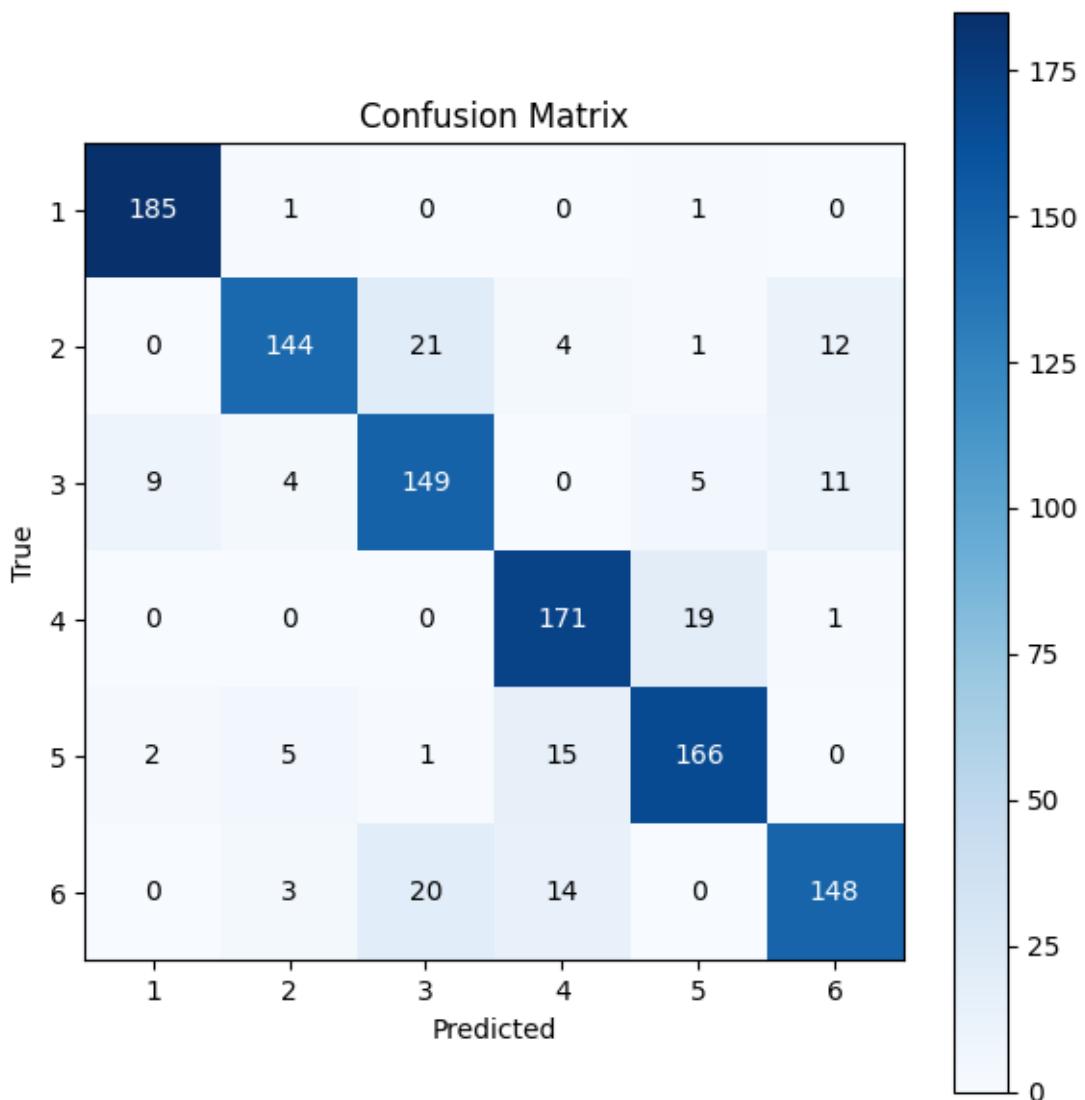
Tuning RF: 78% | 7/9 [02:08<00:36, 18.00s/it]
{'max_depth': 10, 'n_estimators': 50}: train 0.967±0.0000, val 0.826±0.0010

Tuning RF: 89% | 8/9 [02:24<00:17, 17.35s/it]
{'max_depth': 10, 'n_estimators': 100}: train 0.970±0.0000, val 0.830±0.0010

Tuning RF: 100% | 9/9 [02:56<00:00, 19.60s/it]
{'max_depth': 10, 'n_estimators': 200}: train 0.970±0.0000, val 0.833±0.0010

→ Best params = {'max_depth': None, 'n_estimators': 200} (val acc = 0.837)

```



Classification Report:

	precision	recall	f1-score	support
1	0.944	0.989	0.966	187
2	0.917	0.791	0.850	182
3	0.780	0.837	0.808	178
4	0.838	0.895	0.866	191
5	0.865	0.878	0.871	189
6	0.860	0.800	0.829	185
accuracy			0.866	1112
macro avg	0.867	0.865	0.865	1112

weighted avg	0.868	0.866	0.866	1112
--------------	-------	-------	-------	------

2 Support Vector Machine

Tune C and kernel for an SVM, using the same repeated group CV approach.

```
[13]: def tune_svm(
    X_train, y_train, groups,
    param_grid,
    X_test=None, y_test=None,
    n_splits=5, train_size=0.8, random_state=42
):
    """
    Tune SVM over param_grid with GroupShuffleSplit CV,
    retrain on all training data, and (optionally) evaluate on test.
    """
    best_params, best_val = None, -np.inf

    # group-aware splitter
    cv = GroupShuffleSplit(
        n_splits=n_splits,
        train_size=train_size,
        random_state=random_state
    )

    # grid search
    for params in tqdm(ParameterGrid(param_grid), desc="Tuning SVM"):
        clf = SVC(**params, random_state=random_state)
        res = cross_validate(
            clf, X_train, y_train,
            groups=groups, cv=cv,
            return_train_score=True, n_jobs=-1
        )
        tr_m, tr_v = res['train_score'].mean(), res['train_score'].var()
        v_m, v_v = res['test_score'].mean(), res['test_score'].var()
        print(f"{params}: train {tr_m:.3f}±{tr_v:.4f}, val {v_m:.3f}±{v_v:.4f}")

        if v_m > best_val:
            best_val, best_params = v_m, params

    print(f"\n+ Best params = {best_params} (val acc = {best_val:.3f})\n")

    # retrain on full train set
    best_model = SVC(**best_params, random_state=random_state)
    best_model.fit(X_train, y_train)
```

```

# final evaluation & single CM display
if X_test is not None and y_test is not None:
    y_pred = best_model.predict(X_test)
    labels = np.unique(y_test) # e.g. array([1,2,3,4,5,6])
    cm = confusion_matrix(y_test, y_pred, labels=labels)

    # plot heatmap with properly aligned ticks
    fig, ax = plt.subplots(figsize=(6,6))
    im = ax.imshow(cm, interpolation='nearest', cmap='Blues')

    # positions 0..len(labels)-1, but labeled with actual class names
    tick_pos = np.arange(len(labels))
    ax.set_xticks(tick_pos)
    ax.set_xticklabels(labels)
    ax.set_yticks(tick_pos)
    ax.set_yticklabels(labels)

    ax.set_xlabel("Predicted")
    ax.set_ylabel("True")
    ax.set_title("Confusion Matrix")

    # annotate counts
    thresh = cm.max() / 2
    for i in range(len(labels)):
        for j in range(len(labels)):
            color = "white" if cm[i, j] > thresh else "black"
            ax.text(j, i, cm[i, j], ha="center", va="center", color=color)

    fig.colorbar(im, ax=ax)
    plt.tight_layout()
    plt.show()

    # classification report
    print("\nClassification Report:\n",
          classification_report(y_test, y_pred, digits=3))

return best_model

```

[14]:

```

svm_model = tune_svm(X_train=Xtr_p,y_train=ytr_ft, groups=grp_tr, param_grid={‘C’: [0.1, 1, 10], ‘kernel’: [‘linear’, ‘rbf’], ‘gamma’: [‘scale’, ‘auto’]}, X_test=Xte_p,
y_test=yte_ft )

```

Tuning SVM: 8% | 1/12 [00:04<00:47, 4.31s/it]

{'C': 0.1, 'gamma': 'scale', 'kernel': 'linear'}: train 0.883±0.0000, val 0.826±0.0005

```

Tuning SVM: 17% | 2/12 [00:22<02:04, 12.43s/it]
{'C': 0.1, 'gamma': 'scale', 'kernel': 'rbf'}: train 0.864±0.0000, val
0.824±0.0002

Tuning SVM: 25% | 3/12 [00:29<01:28, 9.85s/it]
{'C': 0.1, 'gamma': 'auto', 'kernel': 'linear'}: train 0.883±0.0000, val
0.826±0.0005

Tuning SVM: 33% | 4/12 [00:49<01:51, 13.93s/it]
{'C': 0.1, 'gamma': 'auto', 'kernel': 'rbf'}: train 0.857±0.0000, val
0.788±0.0004

Tuning SVM: 42% | 5/12 [00:55<01:18, 11.18s/it]
{'C': 1, 'gamma': 'scale', 'kernel': 'linear'}: train 0.889±0.0000, val
0.829±0.0005

Tuning SVM: 50% | 6/12 [01:11<01:16, 12.83s/it]
{'C': 1, 'gamma': 'scale', 'kernel': 'rbf'}: train 0.949±0.0000, val
0.860±0.0004

Tuning SVM: 58% | 7/12 [01:19<00:55, 11.03s/it]
{'C': 1, 'gamma': 'auto', 'kernel': 'linear'}: train 0.889±0.0000, val
0.829±0.0005

Tuning SVM: 67% | 8/12 [01:33<00:48, 12.16s/it]
{'C': 1, 'gamma': 'auto', 'kernel': 'rbf'}: train 0.975±0.0000, val 0.851±0.0003

Tuning SVM: 75% | 9/12 [02:04<00:53, 17.86s/it]
{'C': 10, 'gamma': 'scale', 'kernel': 'linear'}: train 0.890±0.0000, val
0.828±0.0005

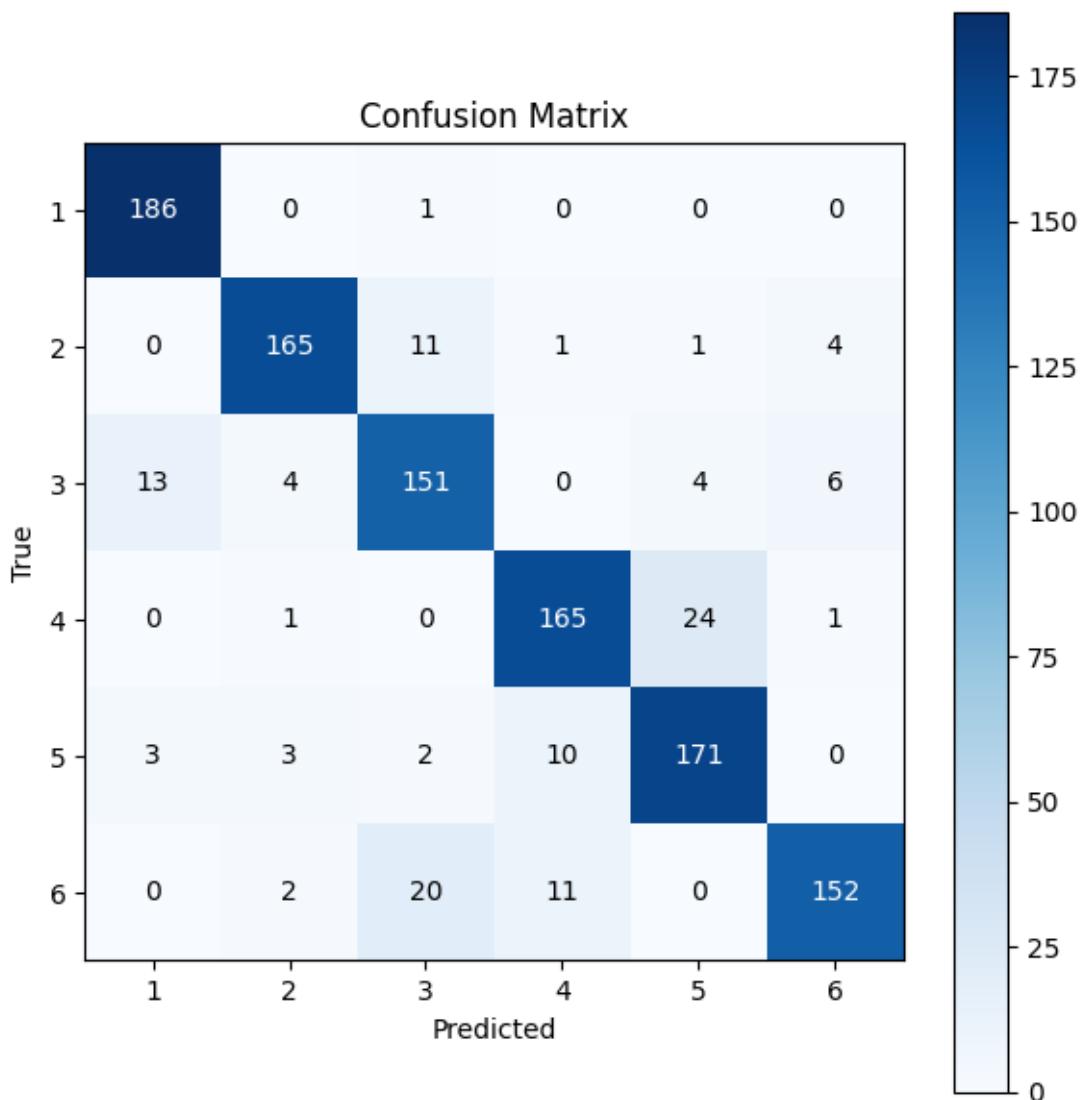
Tuning SVM: 83% | 10/12 [02:16<00:32, 16.07s/it]
{'C': 10, 'gamma': 'scale', 'kernel': 'rbf'}: train 0.993±0.0000, val
0.859±0.0002

Tuning SVM: 92% | 11/12 [02:46<00:20, 20.43s/it]
{'C': 10, 'gamma': 'auto', 'kernel': 'linear'}: train 0.890±0.0000, val
0.828±0.0005

Tuning SVM: 100% | 12/12 [03:01<00:00, 15.13s/it]
{'C': 10, 'gamma': 'auto', 'kernel': 'rbf'}: train 0.999±0.0000, val
0.851±0.0003

→ Best params = {'C': 1, 'gamma': 'scale', 'kernel': 'rbf'} (val acc = 0.860)

```



Classification Report:

	precision	recall	f1-score	support
1	0.921	0.995	0.956	187
2	0.943	0.907	0.924	182
3	0.816	0.848	0.832	178
4	0.882	0.864	0.873	191
5	0.855	0.905	0.879	189
6	0.933	0.822	0.874	185
accuracy			0.890	1112
macro avg	0.892	0.890	0.890	1112

weighted avg	0.892	0.890	0.890	1112
